

Convolutional Neural Networks

(3-0-0-3)

Introduction to CNN

January 26, 2023

Image classification

Image Classification: A core computer vision task

Input: image



This image is © NIKITA
licensed under CC-BY 2.0

Output: Assign image to one
of a fixed set of categories



cat
bird
deer
dog
truck

For human, its a trivial task!

Image classification

Problem: Semantic Gap



This image by Nikita is
licensed under CC-BY 2.0

What the computer sees

An image is just a big grid of numbers between [0, 255]:

e.g. 800 x 600 x 3
(3 channels RGB)

Computer sees only numbers between 0 to 255 for gray scale and 0 to 255 with three different channels for RGB images

Image classification

Challenges: Viewpoint Variation

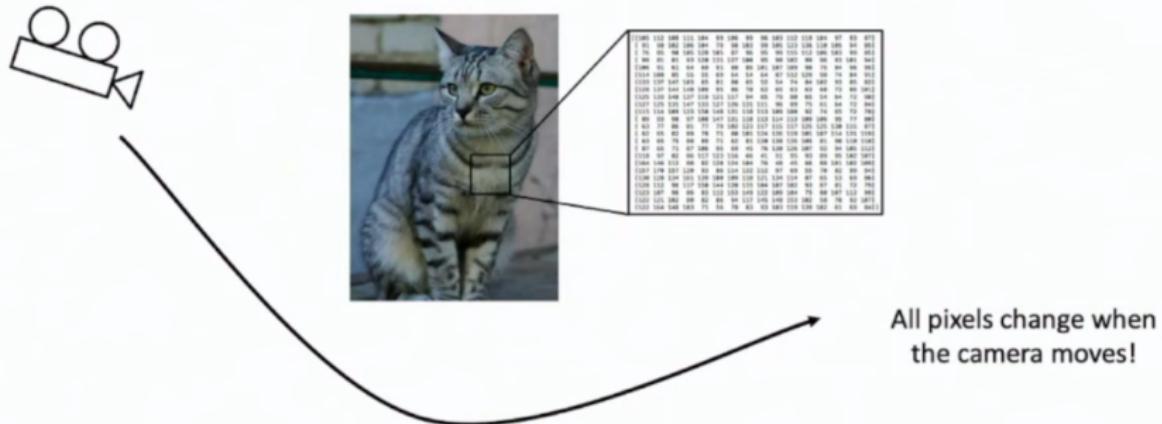


Image classification

Challenges: Intraclass Variation



This image is CC0 1.0 public domain

Image classification

Challenges: Fine-Grained Categories

Maine Coon



[This image](#) is free for use under the [Pixabay License](#).

Ragdoll



[This image](#) - CC0 public domain.

American



[This image](#) - CC0 public domain.

Image classification

Challenges: Background Clutter



[This Image](#) is CC0 1.0 public domain



[This Image](#) is CC0 1.0 public domain

Image classification

Challenges: Illumination Changes



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)



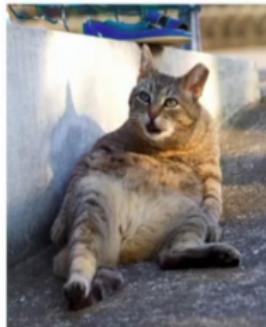
[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)

Image classification

Challenges: Deformation



This image by Umberto Salvagnin is licensed under CC-BY 2.0



This image by Umberto Salvagnin is licensed under CC-BY 2.0



This image by sark bear is licensed under CC-BY 2.0



This image by Tom Thai is licensed under CC-BY 2.0

Image classification

Challenges: Occlusion



This image is CC-BY 3.0 public domain



This image is CC-BY 3.0 public domain



This image is CC-BY 3.0 public domain

Image classification

- Image classification is powerful building block in many computer vision problems
- One cannot hard-code the algorithm with pixels by making indefinite number of rules

Image classification

- Image classification is powerful building block in many computer vision problems
- One cannot hard-code the algorithm with pixles by making indefinite number of rules
- One can try edge detection and hard code the algorithm based on edges (like ears, nose etc)

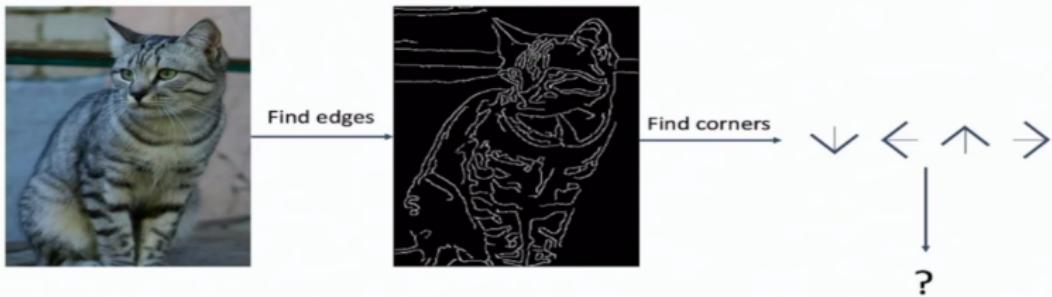


Image classification

- Image classification is powerful building block in many computer vision problems
- One cannot hard-code the algorithm with pixles by making indefinite number of rules
- One can try edge detection and hard code the algorithm based on edges (like ears, nose etc)
- Edge detector may fail based on orientation!!

Image classification

- Image classification is powerful building block in many computer vision problems
- One cannot hard-code the algorithm with pixles by making indefinite number of rules
- One can try edge detection and hard code the algorithm based on edges (like ears, nose etc)
- Edge detector may fail based on orientation!!
- For different problem like identifying car, this algorithm you have developed may not work!

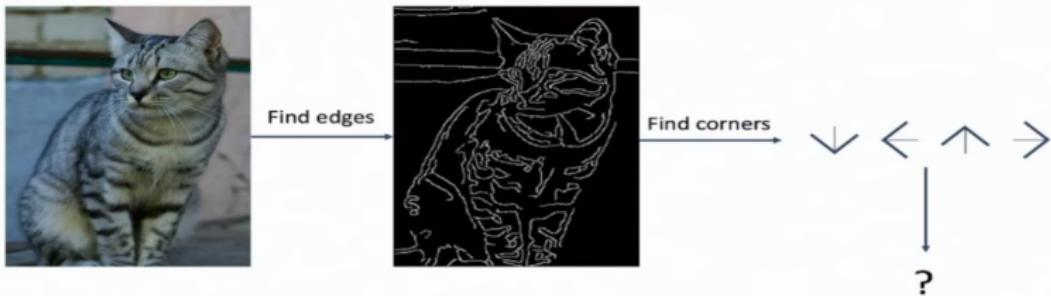
Image classification

- Image classification is powerful building block in many computer vision problems
- One cannot hard-code the algorithm with pixles by making indefinite number of rules
- One can try edge detection and hard code the algorithm based on edges (like ears, nose etc)
- Edge detector may fail based on orientation!!
- For different problem like identifying car, this algorithm you have developed may not work!

So we want some generalised approach where we dont require human knowledge!!

Image classification

- Image classification is powerful building block in many computer vision problems
- One cannot hard-code the algorithm with pixles by making indefinite number of rules
- One can try edge detection and hard code the algorithm based on edges (like ears, nose etc)



Convolutional Neural Networks

- CNN are specialised architectures which works well with visual data (i.e) images and videos
- CNN is largely responsible for revolutionizing 'deep learning' by setting new benchmarks for many image processing tasks!
- CNN uses many of the working principles of animal visual system [\[Read more here\]](#)
- Rapid inventions in computer vision
 - Image classification
 - Object detection
 - Neural style transfer

Convolutional Neural Networks

- In case of images, the input to the neural network becomes very large

Consider an RYB image with array size of 64×64

$$64 \times 64 \times 3 \Rightarrow 12288$$

$$1000 \times 1000 \times 3 \Rightarrow 3 \times 10^6$$

The input to the layer (X) is 3 million values

Convolutional Neural Networks

- In case of images, the input to the neural network becomes very large

Consider an RYB image with array size of 64×64

$$64 \times 64 \times 3 \Rightarrow 12288$$

$$1000 \times 1000 \times 3 \Rightarrow 3 \times 10^6$$

The input to the layer (X) is 3 million values

- If we have first hidden layer of 1000 units, then the order of trainable parameters

$$W^{[1]} \Rightarrow 1000 \times 3 \text{ million}$$

This matrix is very large, NN may overfit! The computational requirement is not feasible!

Convolutional using Edge detection

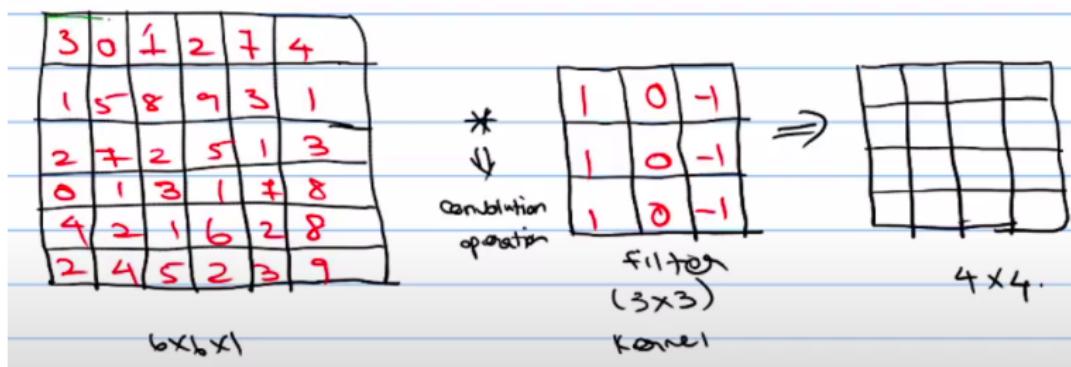
- Similar to cameras where we add filters to have some special effects, here we use array of no's as filters to extract special effects from the image!
- Consider an gray scale image with array size 6×6

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

$6 \times 6 \times 1$

Convolutional using Edge detection

- Similar to cameras where we add filters to have some special effects, here we use array of no's as filters to extract special effects from the image!
- Consider a gray scale image with array size 6×6



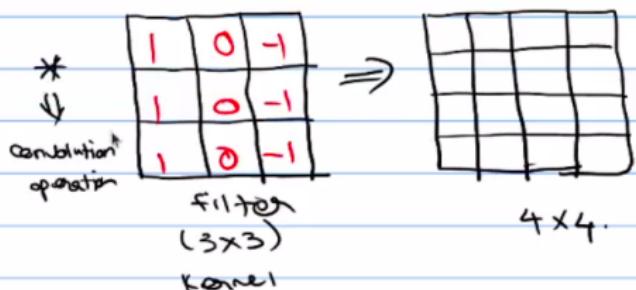
- Filter is also called as “Kernel”

Convolutional using Edge detection

- Similar to cameras where we add filters to have some special effects, here we use array of no's as filters to extract special effects from the image!
- Consider an gray scale image with array size 6×6

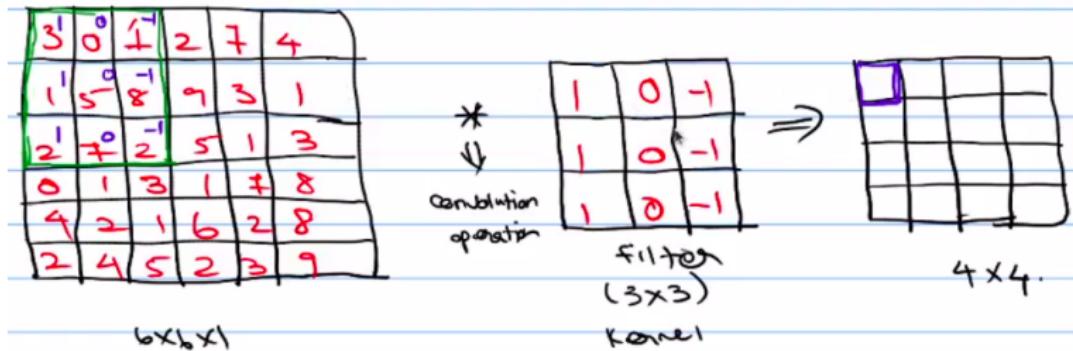
3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

$6 \times 6 \times 1$



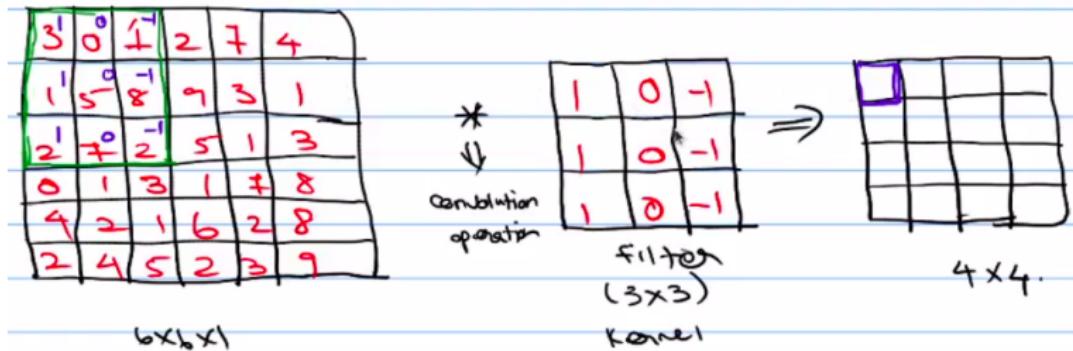
Convolutional using Edge detection

- Similar to cameras where we add filters to have some special effects, here we use array of no's as filters to extract special effects from the image!
- Consider an gray scale image with array size 6×6



Convolutional using Edge detection

- Similar to cameras where we add filters to have some special effects, here we use array of no's as filters to extract special effects from the image!
- Consider an gray scale image with array size 6×6



Position 1:

$$= (3*1) + (0*0) + (1*-1) + (1*1) + (5*0) + (8*-1) + (2*1) + (7*0) + (2*-1) = -5$$

Convolutional using Edge detection

3	0	1	2	7	4
1	5	8	9	3	1
2	7	0	-1	5	1
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

$6 \times 6 \times 1$

$*$
 \downarrow
 Convolution
 operation

filter
 (3×3)

kernel

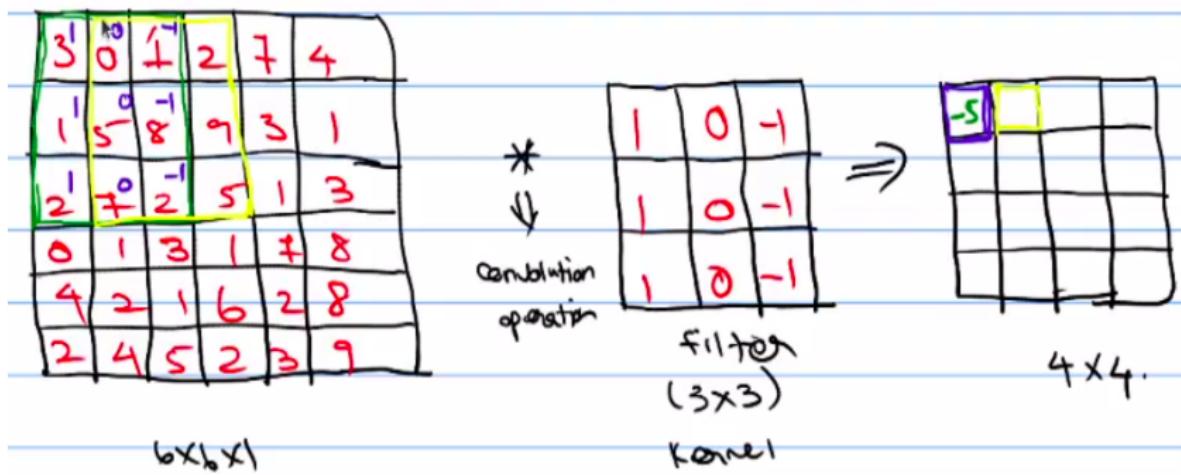
1	0	-1
1	0	-1
1	0	-1



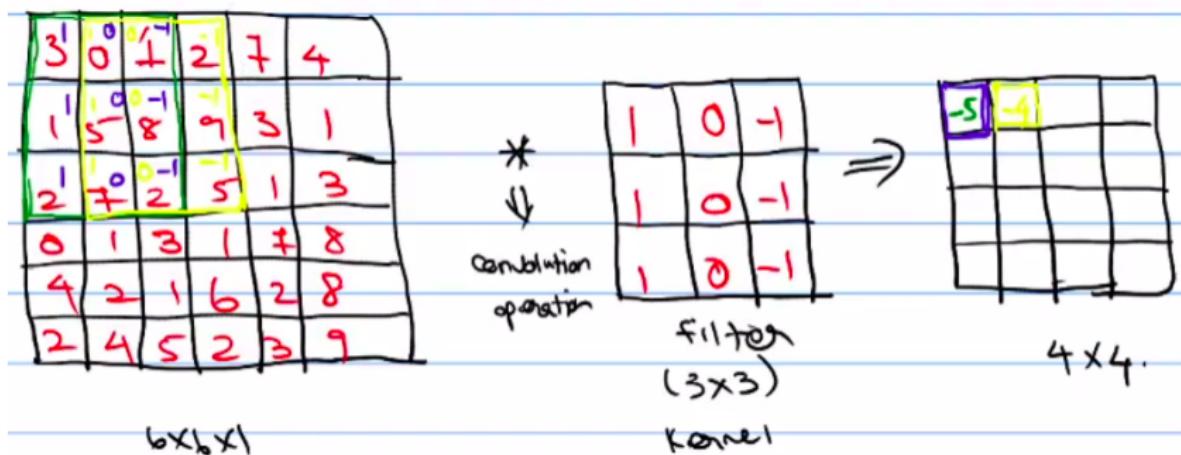
-5			

4×4

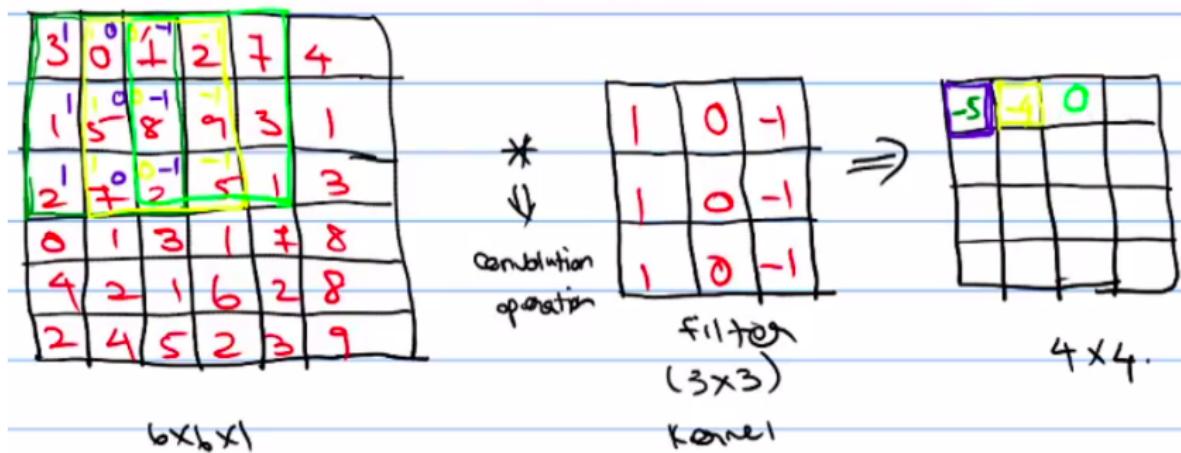
Convolutional using Edge detection



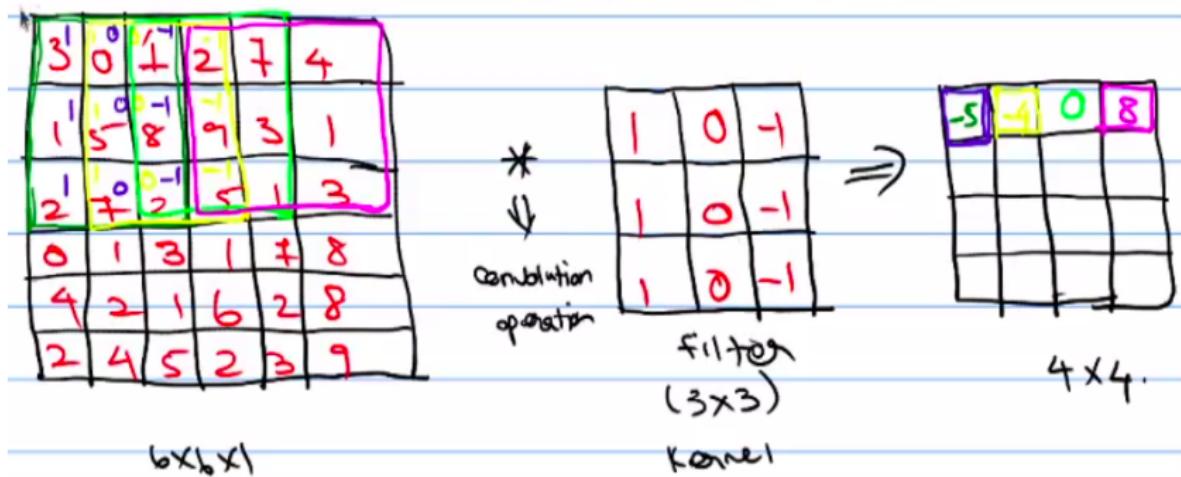
Convolutional using Edge detection



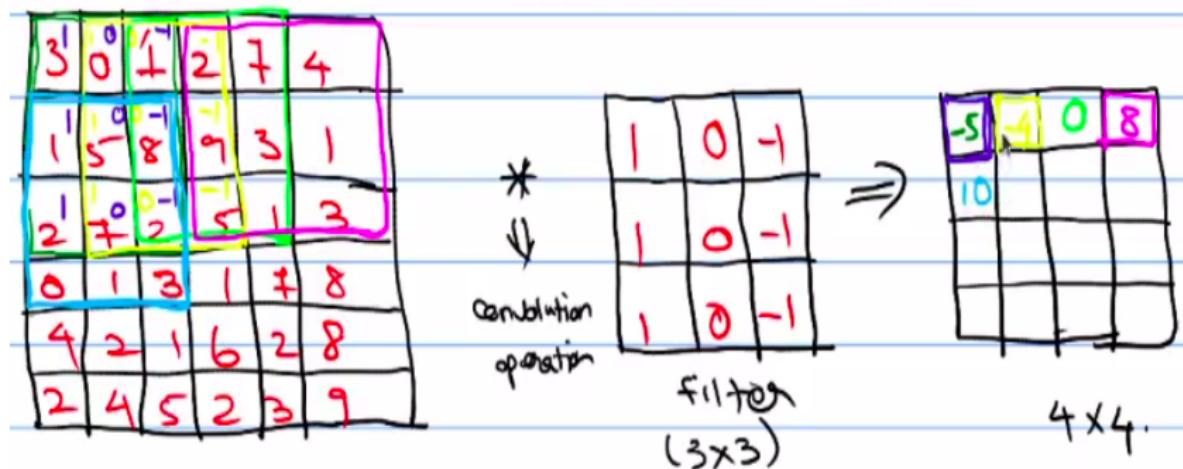
Convolutional using Edge detection



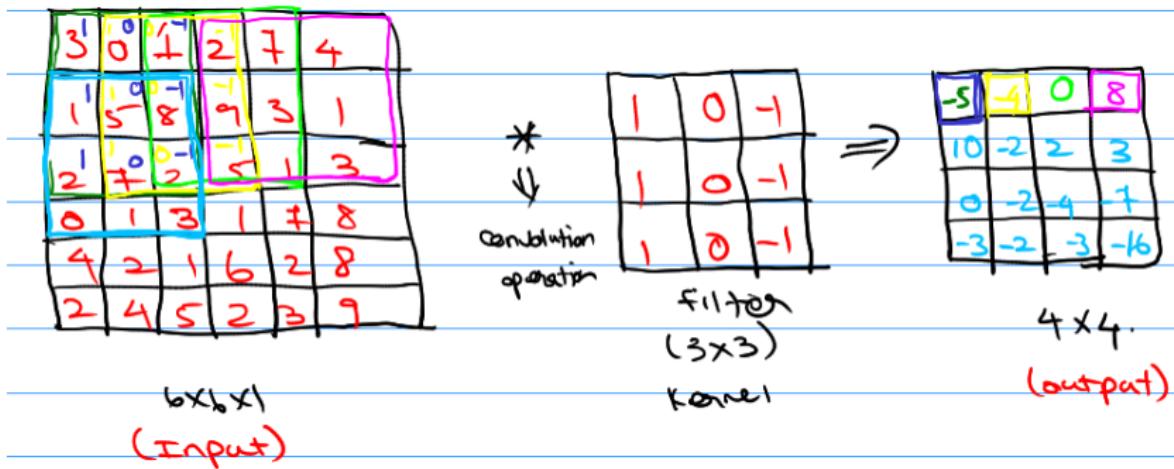
Convolutional using Edge detection



Convolutional using Edge detection



Convolutional using Edge detection



- This type of filter detects “Vertical edges” and hence called as **Vertical edge detection filter**
- The generalised way of calculating output size
$$(n \times n) * (f \times f) \Rightarrow (n - f + 1) * (n - f + 1)$$

Vertical Edge detection

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

*

1	0	-1
1	0	-1
1	0	-1



=



An example of a vertical edge detection. Here we have light to dark transition.

- Vertical edge detection filter is used

Vertical Edge detection

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

*

1	0	-1
1	0	-1
1	0	-1



=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0



An example of a vertical edge detection. Here we have light to dark transition.

- White in the output indicates transition from bright to dark

Vertical Edge detection

0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10

*

1	0	-1
1	0	-1
1	0	-1



=

0	-30	-30	0
0	-30	-30	0
0	-30	-30	0
0	-30	-30	0



An example of vertical edge detection of dark to light transition

- Black in the output indicates transition from dark to bright

Horizontal Edge detection

255	255	255	255	255	255
255	255	255	255	255	255
255	255	255	255	255	255
255	255	255	255	255	255
10	10	10	10	10	10
10	10	10	10	10	10
10	10	10	10	10	10
10	10	10	10	10	10

8 X 6

*

1	1	1
0	0	0
-1	-1	-1

3 X 3

6 X 4

=

0	0	0	0
0	0	0	0
735	735	735	735
735	735	735	735
0	0	0	0
0	0	0	0



- White in the output indicates transition from bright to dark

Edge detection

Input Image						
3	1	6	2	8	0	
4	8	5	9	1	5	
6	8	0	0	1	9	
3	6	9	14	02	-11	
5	2	5	10	07	-11	
8	3	0	15	03	-11	

6x6

→ convolutional operation

$*$

filter	=	Output
$\begin{matrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{matrix}$ 3x3		$\begin{matrix} 2 & 4 & 1 & -3 \\ -1 & 9 & 10 & -2 \\ 0 & 12 & 4 & -7 \\ 2 & 2 & 2 & 6 \end{matrix}$

$= 1*4 + 1*0 + 1*5 + 0*2 + 0*7 + 0*3 + (-1)*1 + (-1)*1 + (-1)*1 = 6$



Visual for input matrix



Current visual for output matrix

Convolutional Operation

Other filters in Image processing

$$\begin{matrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{matrix}$$

\Rightarrow SOBEL FILTER

$$\begin{matrix} 3 & 0 & -3 \\ 10 & 0 & -10 \\ 3 & 0 & -3 \end{matrix}$$

\Rightarrow SCHWASS
FILTER

- These filter give weightage more to central values

Other filters in Image processing

- Inorder to identify edges in the complex image, rather than using pre-defined filters, one can define them as parameters in NN and obtain them using backpropagation



- Rather than learning only horizontal and vertical edges, this filter may findout angles 45° , 37° , etc..

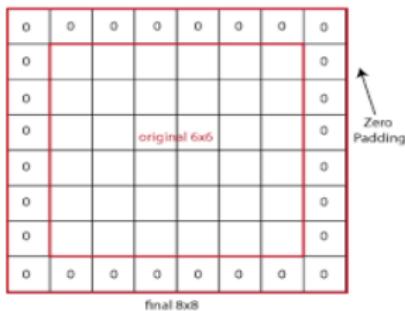
Padding

- When we apply convolution operator everytime, image shrinks!
- Pixels at the corner of the matrix is used only in one of the matrix computation
- Pixels at the center is used in more matrix computation.
- Pixels at the corner are least used

To overcome this problem, we use “**padding**”

Padding

Padding is the process of adding border with “0” to the given image



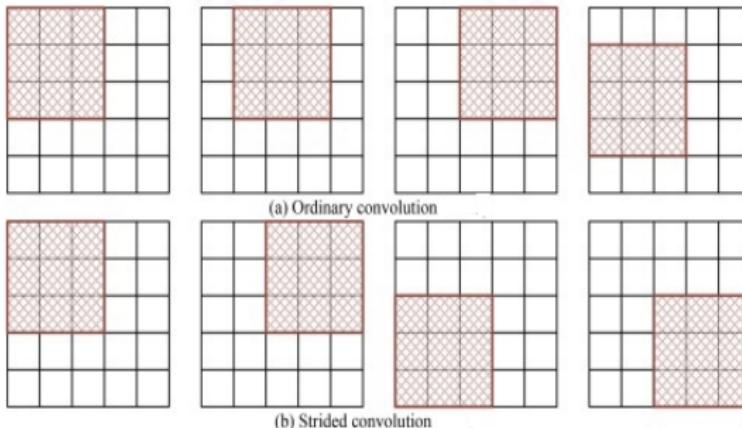
Original image $\Rightarrow 6 \times 6$
With padded image $\Rightarrow 8 \times 8$
Filter $\Rightarrow 3 \times 3$

Output without padding $\Rightarrow 4 \times 4$
Output with padding $\Rightarrow 6 \times 6$

- The general expression for output size with padding is $(n \times n) * (f \times f) \Rightarrow (n + 2p - f + 1) * (n + 2p - f + 1)$
- In padding, sometimes instead of zero, one use the value of edge pixels

Strided Convolution

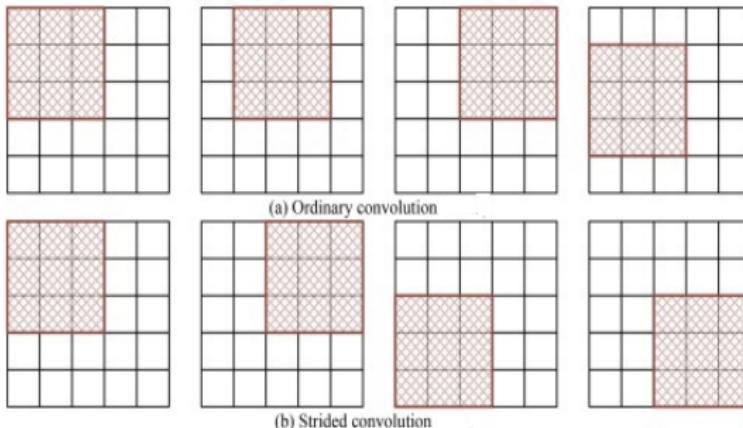
Rather than moving to the next cell continuously, we jump based on stride number



- Row-1 \Rightarrow Ordinary convolution (Stride = 1)
- Row-2 \Rightarrow Strided convolution (Stride=2)
- Stride (jump) is on both horizontal and vertical side

Strided Convolution

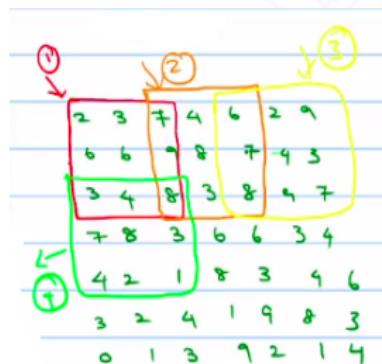
Rather than moving to the next cell continuously, we jump based on stride number



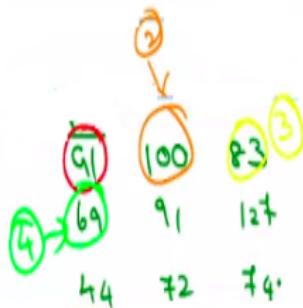
- Row-1 \Rightarrow Ordinary convolution (Stride = 1)
- Row-2 \Rightarrow Strided convolution (Stride=2)
- Stride (jump) is on both horizontal and vertical side

Strided Convolution

Rather than moving to the next cell continuously, we jump based on stride number



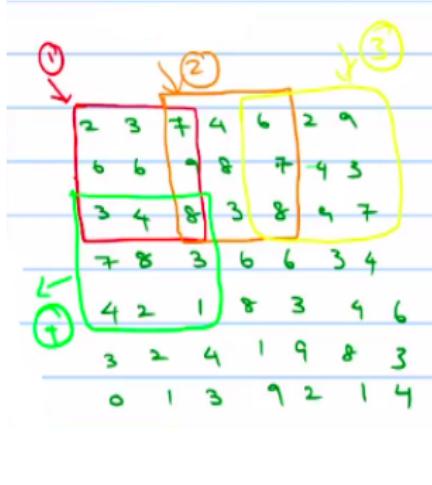
Output:



- The general expression for output size with padding is $(n \times n) * (f \times f) \Rightarrow (\frac{n+2p-f}{S} + 1) * (\frac{n+2p-f}{S} + 1)$ where 'p' is the padding and 'S' is the stride

Strided Convolution

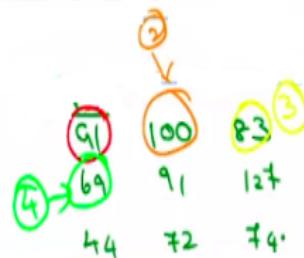
Rather than moving to the next cell continuously, we jump based on stride number



Filter:

$$\begin{matrix} 3 & 4 & 4 \\ 1 & 0 & 2 \\ -1 & 0 & 3 \end{matrix}$$

Output:

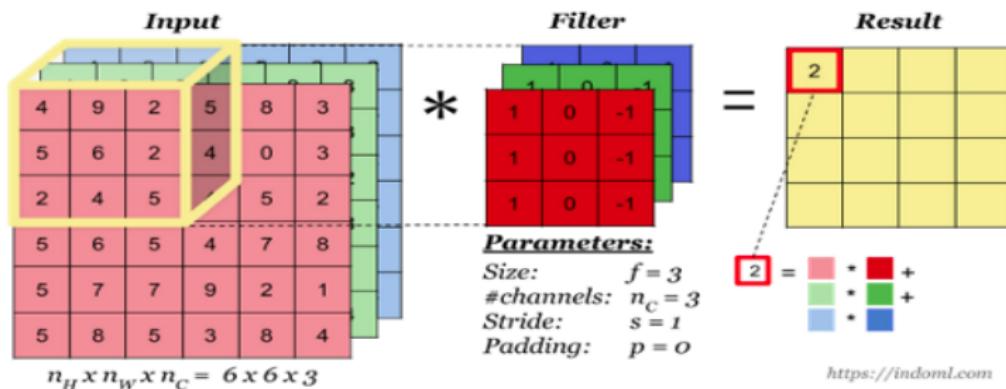


Strided convolution

- The rate at which we are scanning the image is called stride (rate at which we move the filter)
- When we use higher stride, if the length is not sufficient for convolution, we will pad at the edges
- If one is interested in very fine features, then stride should be small
- If one is interested in macro features then stride should be large

Convolution over volume

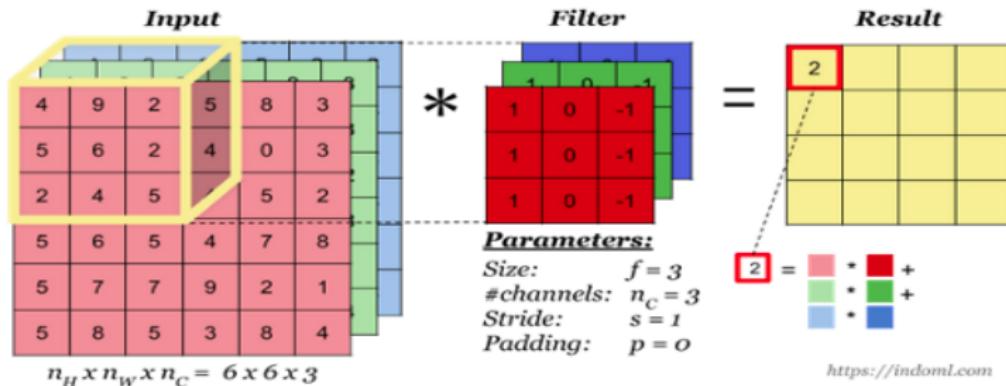
- Till now, we have discussed convolution over 2D images (Gray scale). Now, we will discuss about convolution over volume (more than 2D) for RGB images



<https://indoml.com>

Convolution over volume: element-wise product in each channel then adding them together.
 Source: <https://indoml.com/2018/03/07/student-notes-convolutional-neural-networks-cnn-introduction/>

Convolution over volume

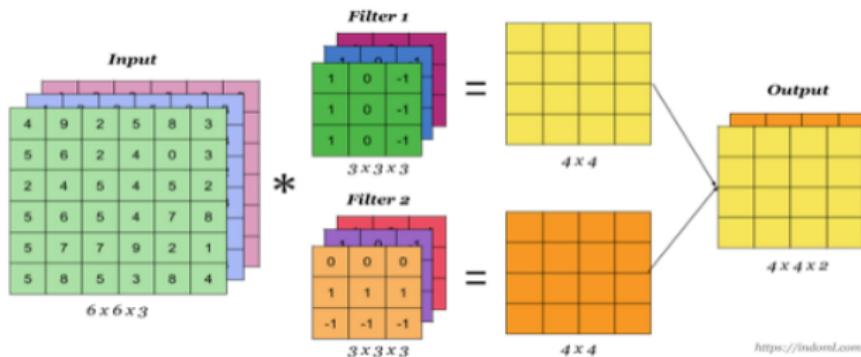


Convolution over volume: element-wise product in each channel then adding them together.

Source: <https://indoml.com/2018/03/07/student-notes-convolutional-neural-networks-cnn-introduction/>

- Each of 27 number in each filter channel is multiplied with 27 number in images and add everything to compute the values
- Number of channels in the image should match with number of channels in the filter

Convolution with multiple filters



Convolution using 2 filters. The output is two 4×4 feature maps. One from each filter
 Source:
<https://indoml.com/2018/03/07/student-notes-convolutional-neural-networks-cnn-introduction/>

- While using multiple filters, output is stacked!

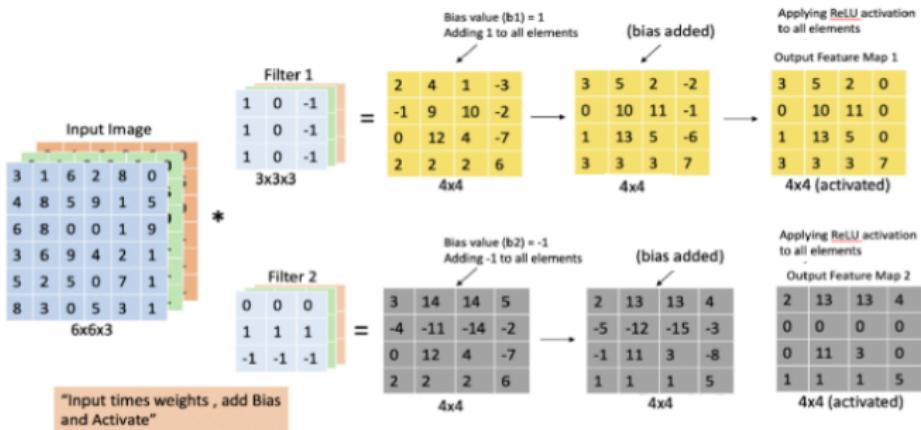
$\text{Image} \Rightarrow n \times n \times n_c$

$\text{Filter} \Rightarrow f \times f \times n_c$

If there are n'_c number of filters, then

$\text{Output} \Rightarrow (n-f+1) \times (n-f+1) \times n'_c$

Single Convolution layer

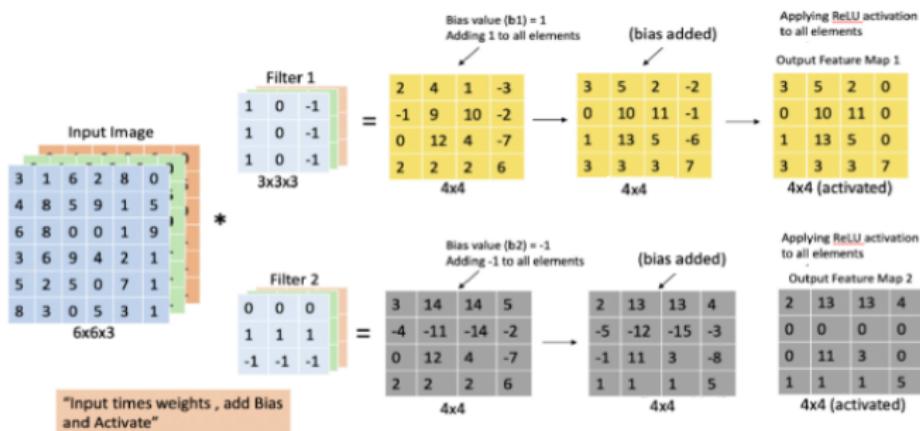


Original image * Filter 1 $\rightarrow \text{relu}[(4 \times 4) + b_1] \rightarrow 4 \times 4$

Original image * Filter 2 $\rightarrow \text{relu}[(4 \times 4) + b_2] \rightarrow 4 \times 4$

$x \cdot w \rightarrow w^T x \rightarrow g(w^T x + b) \rightarrow y$

Single Convolution layer

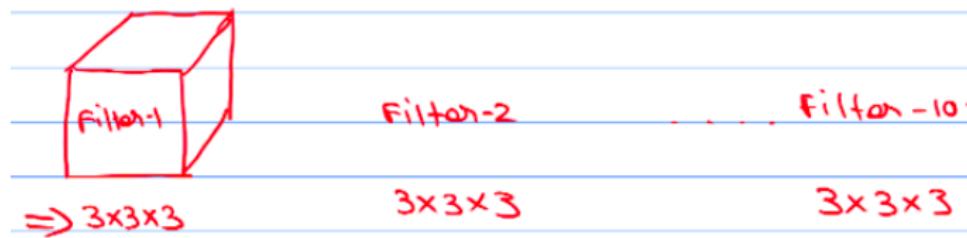


Entire computation of convolutional layer

- This process is very similar to forward propagation
- Filters are acting as weights and it is multiplied with input X
- To this product, add bias and then apply activation
- The output from activation is called as “**feature map**”

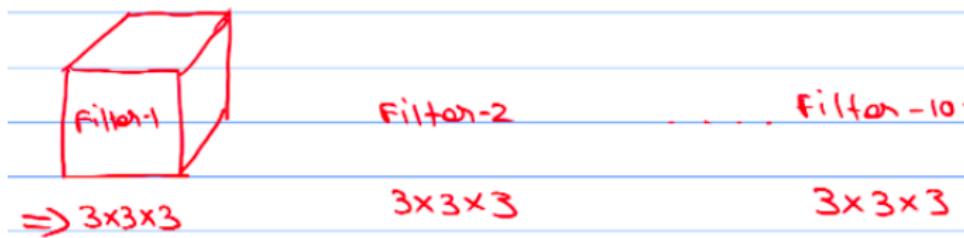
Number of parameters in one layer

- Consider we have 3×3 RYB filters which are 10 in numbers.
- Trainable parameters are W and b. **Calculate total number of parameters.**



Number of parameters in one layer

- Consider we have 3×3 RYB filters which are 10 in numbers.
- Trainable parameters are W and b. **Calculate total number of parameters.**



- For each filter, there will be 27 weights and one bias totalling to 28 parameters
- For 10 filters, total parameters = $28 \times 10 = 280$ parameters
- Irrespective of image size, the number of parameters remain same.

Number of parameters in L^{th} layer

If layer L is a convolutional layer,

$f^{[L]}$ \Rightarrow Filter size

$p^{[L]}$ \Rightarrow Padding

$s^{[L]}$ \Rightarrow Stride

$n_c^{[L]}$ \Rightarrow Number of Filter

$$Input \Rightarrow n_H^{[L-1]} \times n_W^{[L-1]} \times n_c^{[L-1]} \quad Output \Rightarrow n_H^{[L]} \times n_W^{[L]} \times n_c^{[L]}$$

Number of parameters in L^{th} layer

If layer L is a convolutional layer,

$f^{[L]}$ \Rightarrow Filter size

$p^{[L]}$ \Rightarrow Padding

$s^{[L]}$ \Rightarrow Stride

$n_c^{[L]}$ \Rightarrow Number of Filter

$Input \Rightarrow n_H^{[L-1]} \times n_W^{[L-1]} \times n_c^{[L-1]}$ $Output \Rightarrow n_H^{[L]} \times n_W^{[L]} \times n_c^{[L]}$

$n_H^{[L]} \Rightarrow \left(\frac{n^{[L-1]} + 2P^{[L]} - f^{[L]}}{S^{[L]}} + 1 \right)$ $n_W^{[L]} \Rightarrow \left(\frac{n^{[L-1]} + 2P^{[L]} - f^{[L]}}{S^{[L]}} + 1 \right)$

Number of parameters in L^{th} layer

$$Input \Rightarrow n_H^{[L-1]} \times n_W^{[L-1]} \times n_c^{[L-1]}$$

$$Output \Rightarrow n_H^{[L]} \times n_W^{[L]} \times n_c^{[L]}$$

$$n_H^{[L]} \Rightarrow \left(\frac{n^{[L-1]} + 2P^{[L]} - f^{[L]}}{S^{[L]}} + 1 \right)$$

$$n_W^{[L]} \Rightarrow \left(\frac{n^{[L-1]} + 2P^{[L]} - f^{[L]}}{S^{[L]}} + 1 \right)$$

Number of parameters in L^{th} layer

$$Input \Rightarrow n_H^{[L-1]} \times n_W^{[L-1]} \times n_c^{[L-1]} \quad Output \Rightarrow n_H^{[L]} \times n_W^{[L]} \times n_c^{[L]}$$

$$n_H^{[L]} \Rightarrow \left(\frac{n^{[L-1]} + 2P^{[L]} - f^{[L]}}{S^{[L]}} + 1 \right) \quad n_W^{[L]} \Rightarrow \left(\frac{n^{[L-1]} + 2P^{[L]} - f^{[L]}}{S^{[L]}} + 1 \right)$$

If there are n_c channels, then the filter size becomes

$$Filter\ size \Rightarrow f^{[L]} \times f^{[L]} \ n_c^{[L-1]}$$

The filter channel must match with input channel

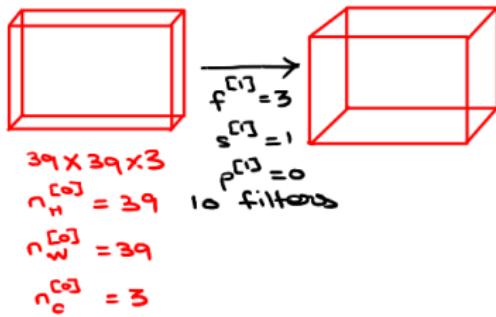
If there are n'_c filters in the L layer

$$Weights \Rightarrow f^{[L]} \times f^{[L]} \times n_c^{[L-1]} \times n_c^{[L]}'$$

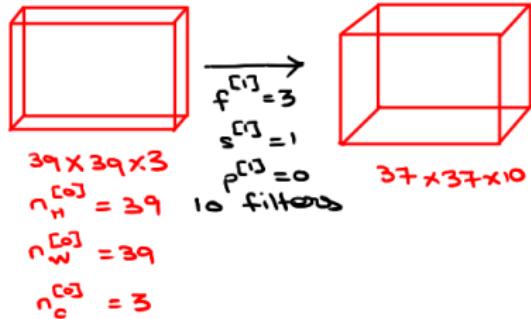
$$Bias \Rightarrow n_c^{[L]}$$

$$Vectorised\ Output\ A^{[L]} \Rightarrow m \times n_H^{[L]} \times n_W^{[L]} \times n_c^{[L]}$$

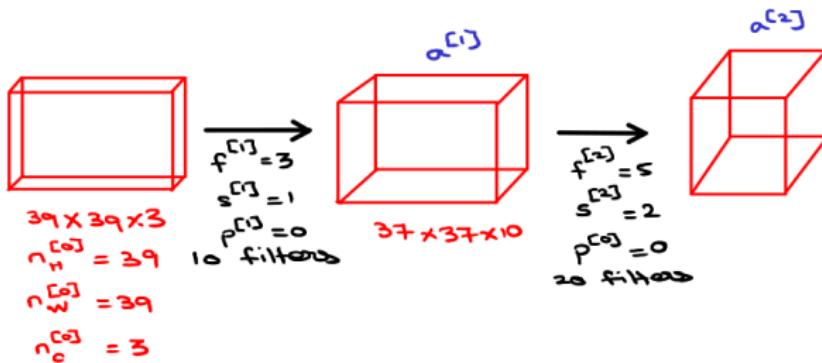
Convolutional Neural Network



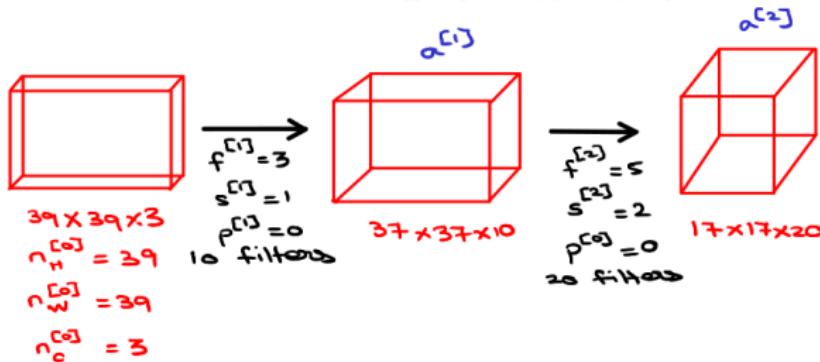
Convolutional Neural Network



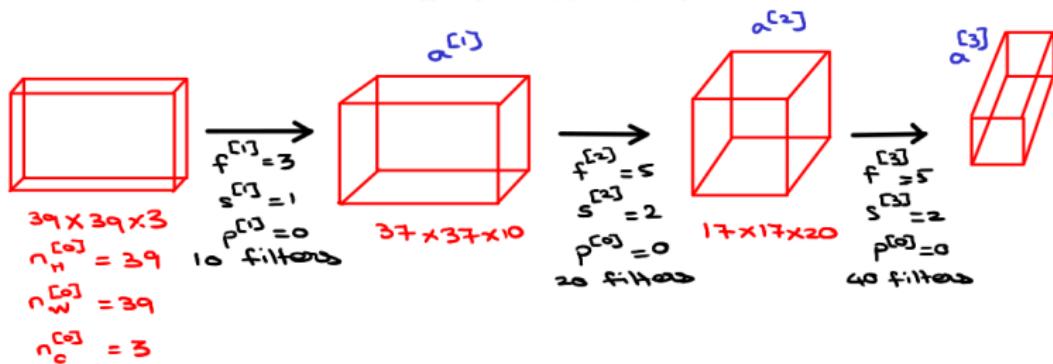
Convolutional Neural Network



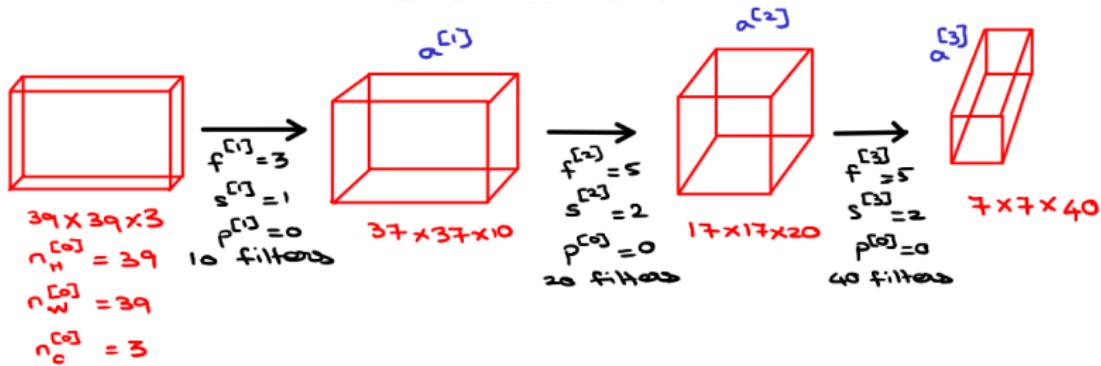
Convolutional Neural Network



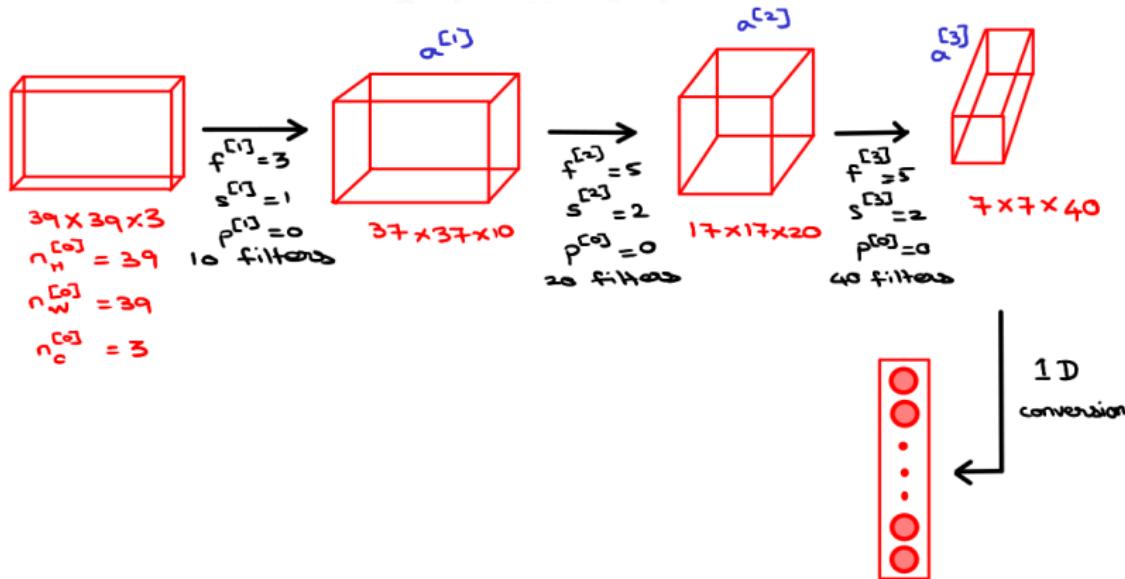
Convolutional Neural Network



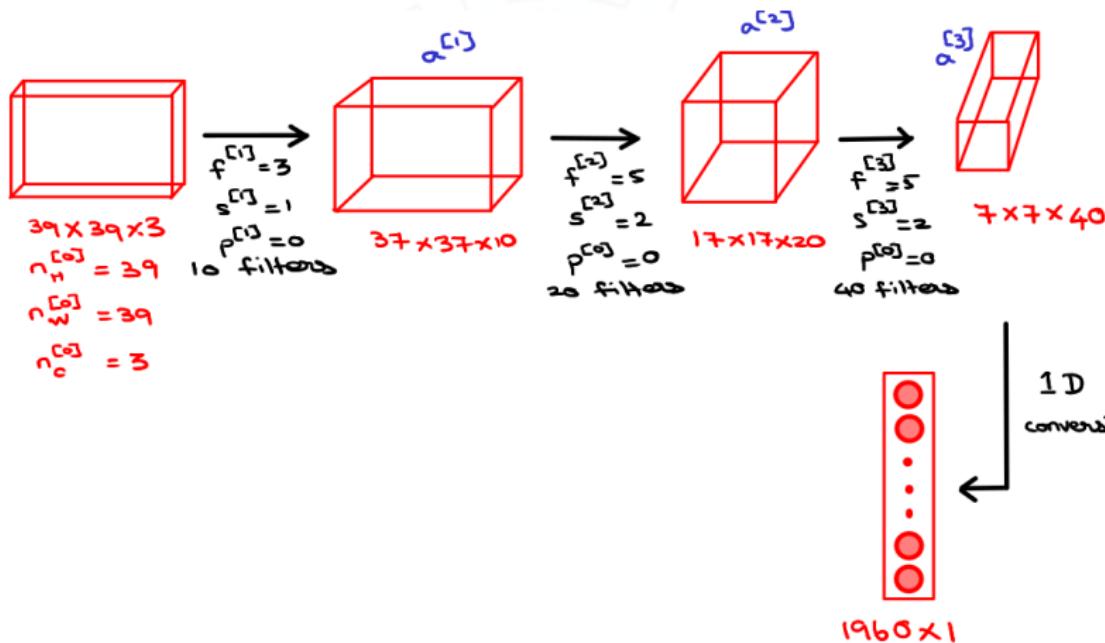
Convolutional Neural Network



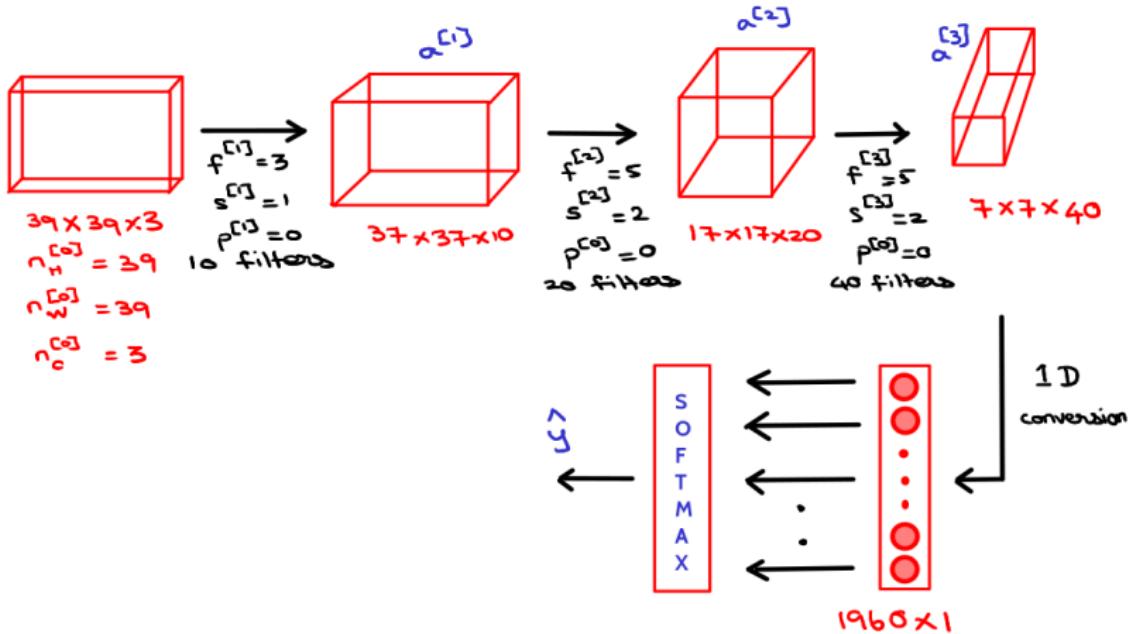
Convolutional Neural Network



Convolutional Neural Network



Convolutional Neural Network



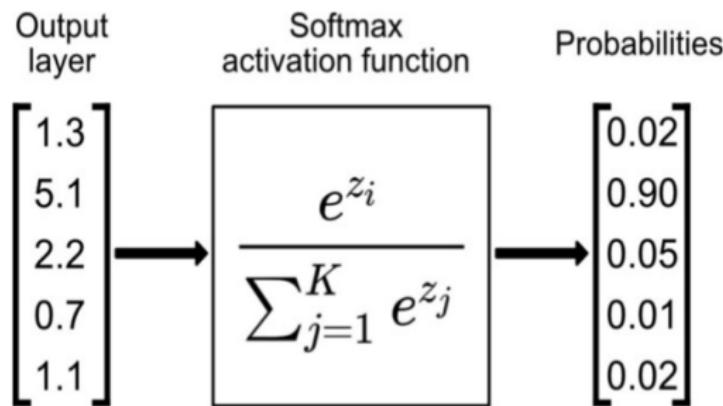
- Image size is going down ($39-37-17-7$)
- Number of filters ($3-10-20-40$)

Convolutional Neural Network: Softmax layer

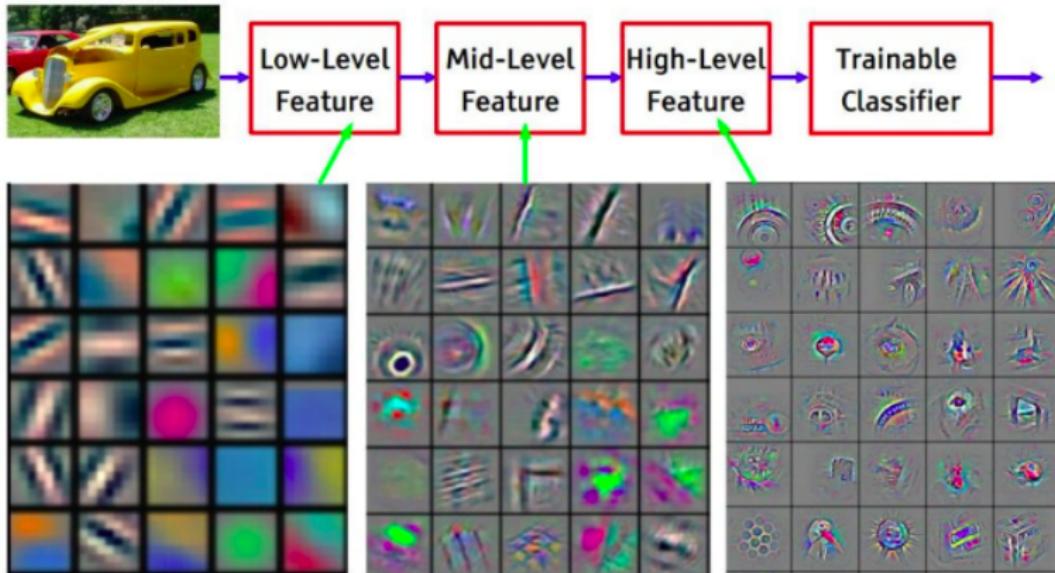
- Used for multi-class classification
- Gives the output probabilities of particular class and adds up to 1
- If there are 'c' classes, then there will be c values
- Array as an input and array as an output

$$\frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

Convolutional Neural Network: Softmax layer



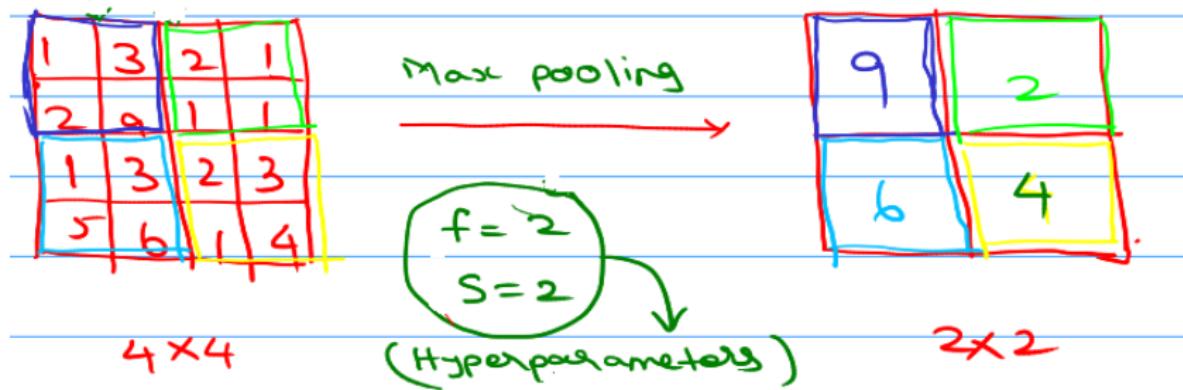
CNN: Features at different levels



Different features recognised at different layers

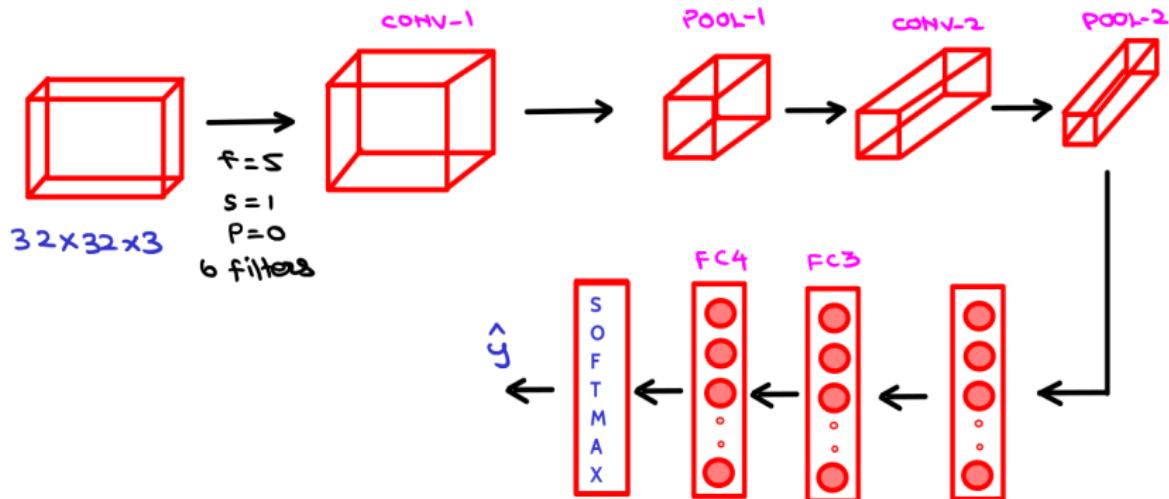
Convolutional Neural Network: Pooling layer

- Pooling layer is used for speedup the computation
- The value of 'f' and 's' are defined by the user
- No learnable parameters in the pooling layer



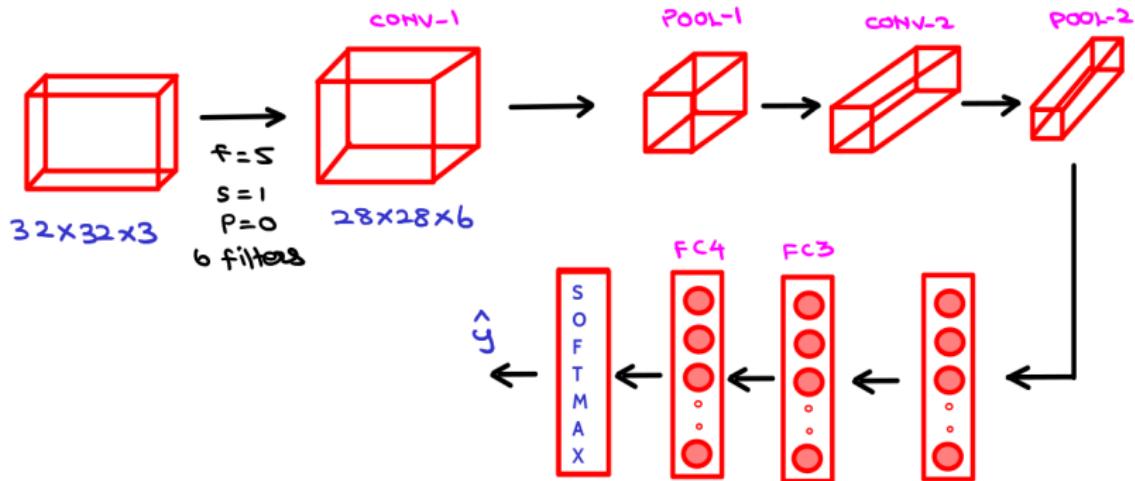
- The generalised output size $\Rightarrow \frac{n+2p-f}{s} + 1$

Parameters in CNN



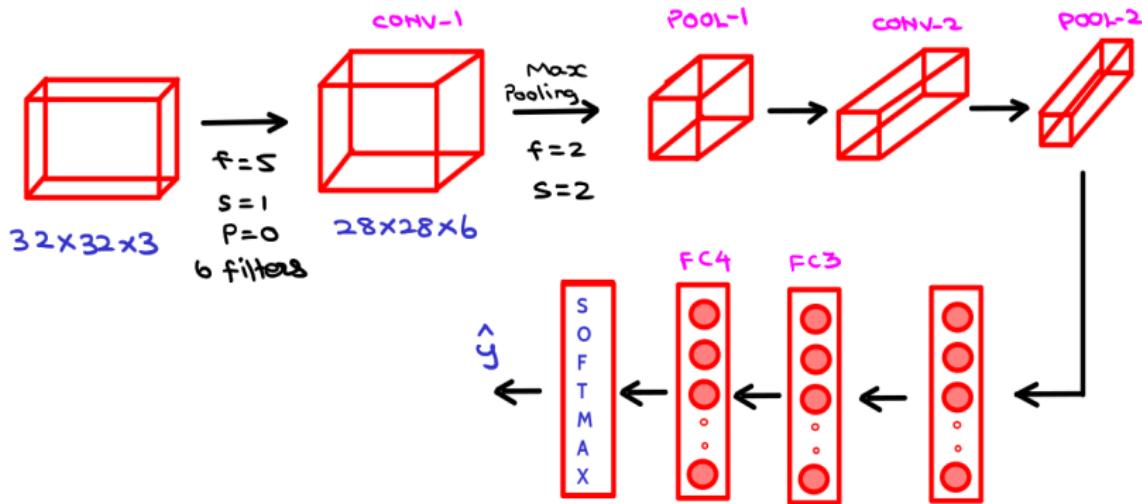
- CONV1-POOL1-CONV2-POOL2-FC3-FC4-Softmax

Parameters in CNN



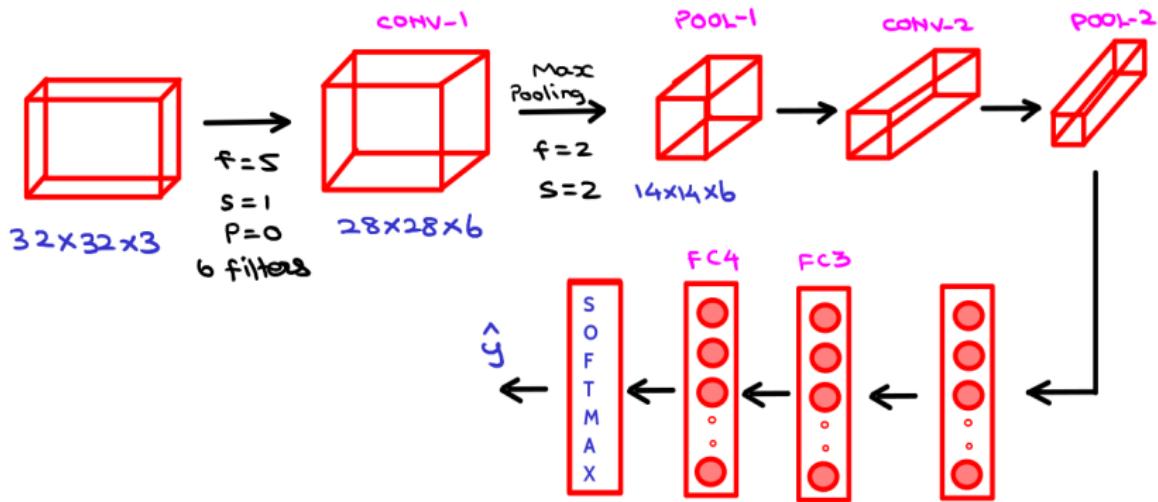
- CONV1-POOL1-CONV2-POOL2-FC3-FC4-Softmax

Parameters in CNN



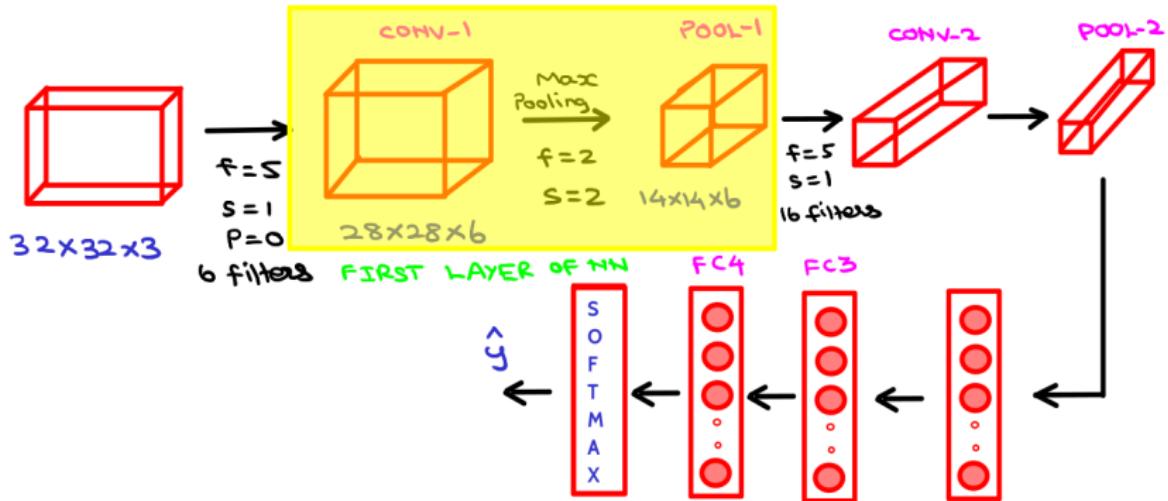
- CONV1-POOL1-CONV2-POOL2-FC3-FC4-Softmax

Parameters in CNN



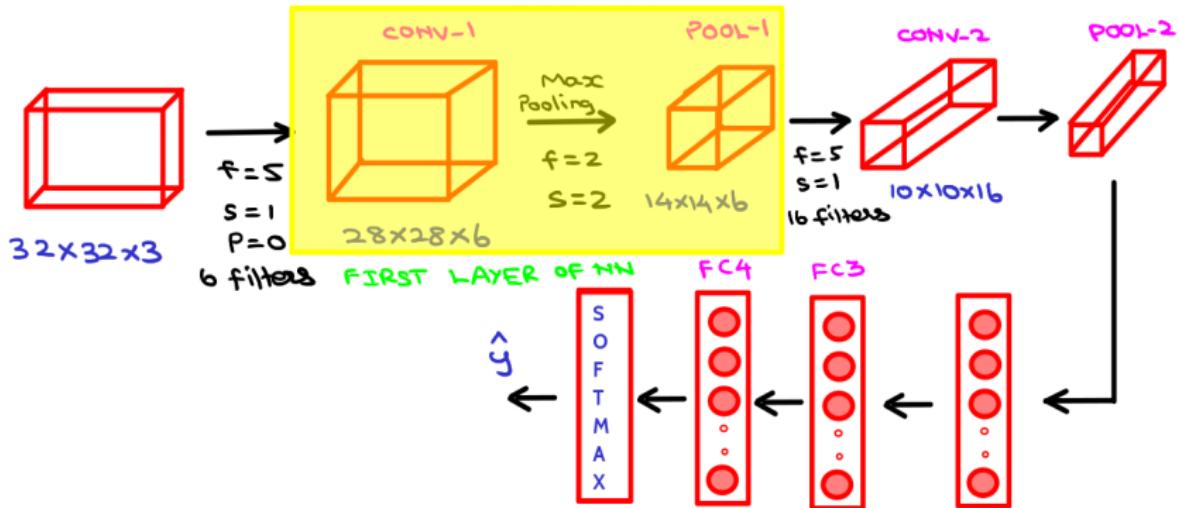
- CONV1-POOL1-CONV2-POOL2-FC3-FC4-Softmax

Parameters in CNN



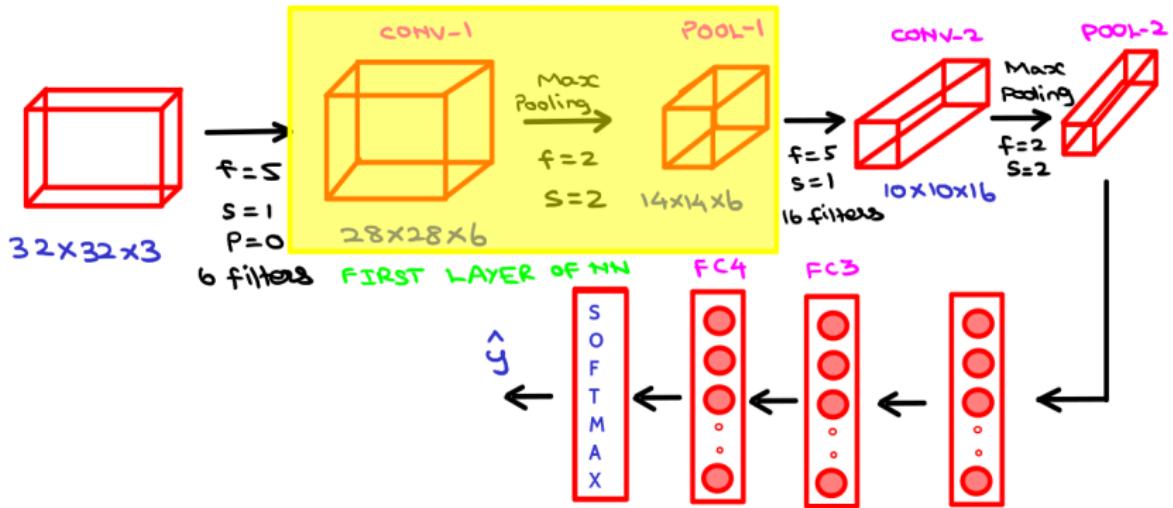
- CONV1-POOL1-CONV2-POOL2-FC3-FC4-Softmax

Parameters in CNN



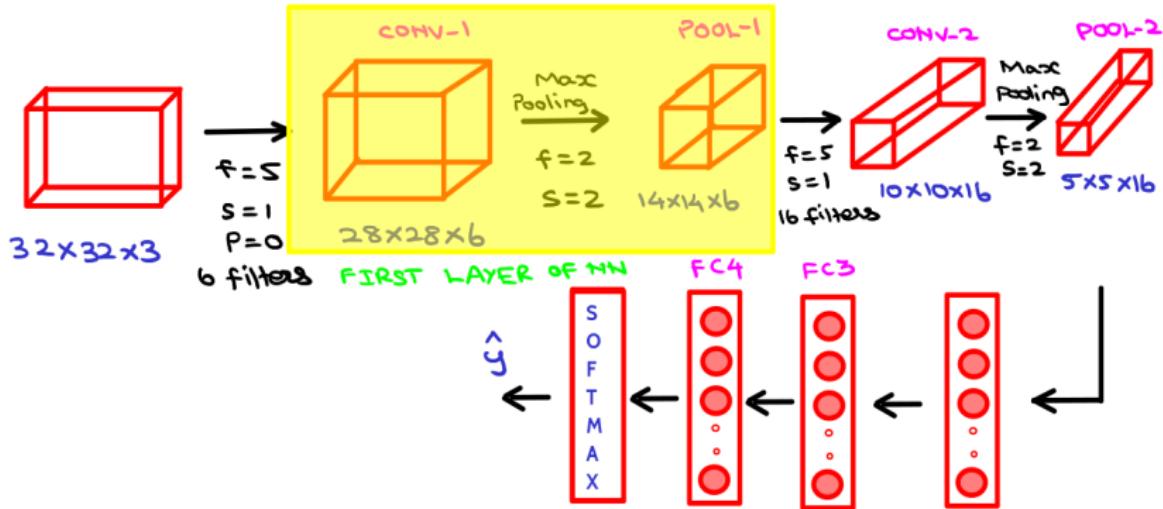
- CONV1-POOL1-CONV2-POOL2-FC3-FC4-Softmax

Parameters in CNN



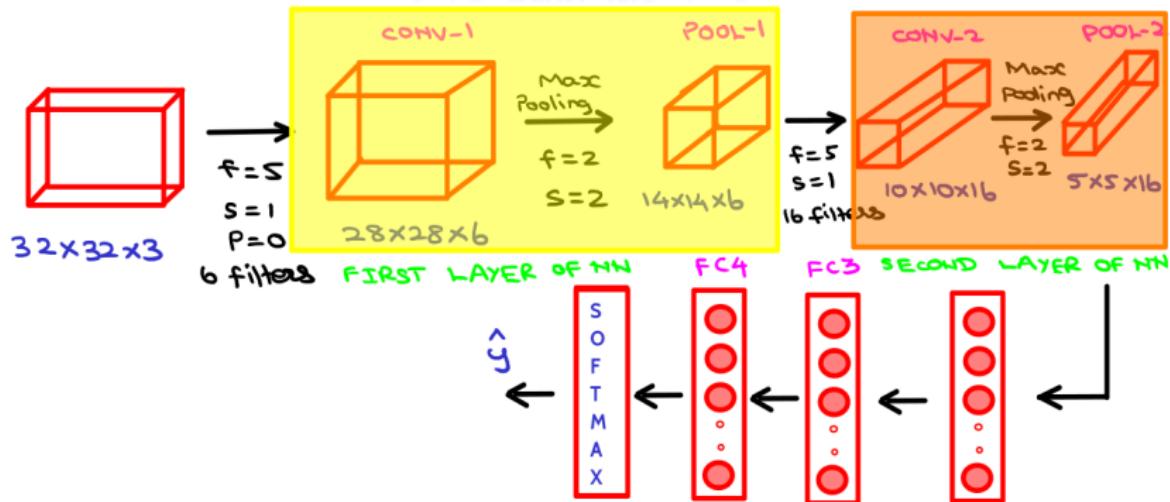
- CONV1-POOL1-CONV2-POOL2-FC3-FC4-Softmax

Parameters in CNN



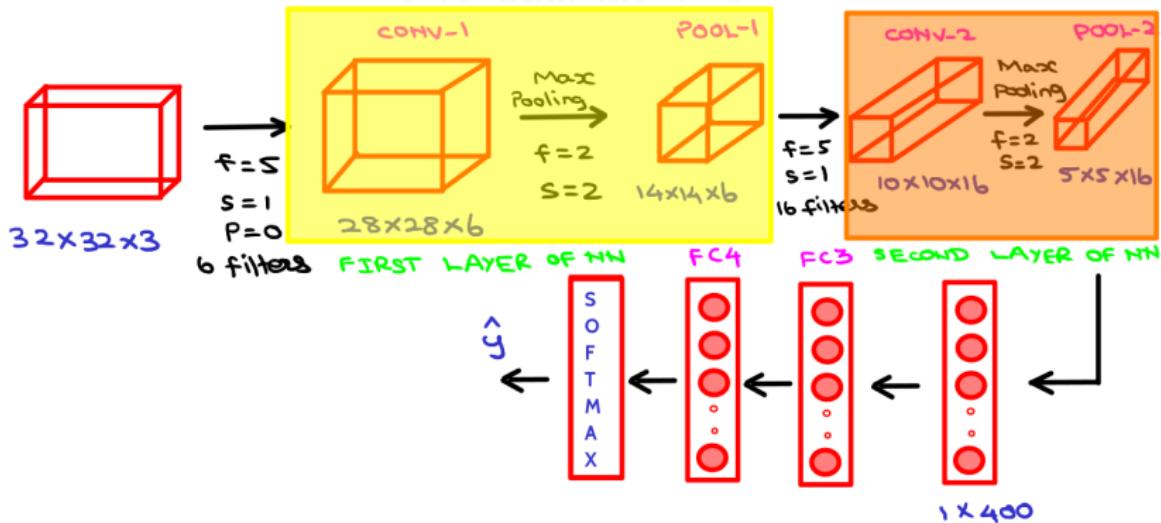
- CONV1-POOL1-CONV2-POOL2-FC3-FC4-Softmax

Parameters in CNN



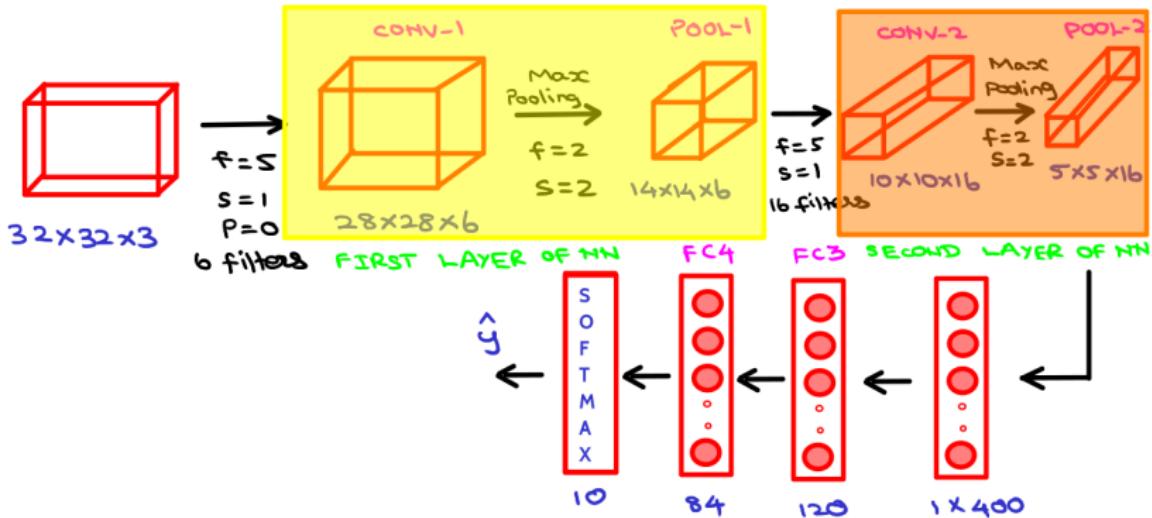
- CONV1-POOL1-CONV2-POOL2-FC3-FC4-Softmax

Parameters in CNN



- CONV1-POOL1-CONV2-POOL2-FC3-FC4-Softmax

Parameters in CNN



- CONV1-POOL1-CONV2-POOL2-FC3-FC4-Softmax

Parameters in CNN

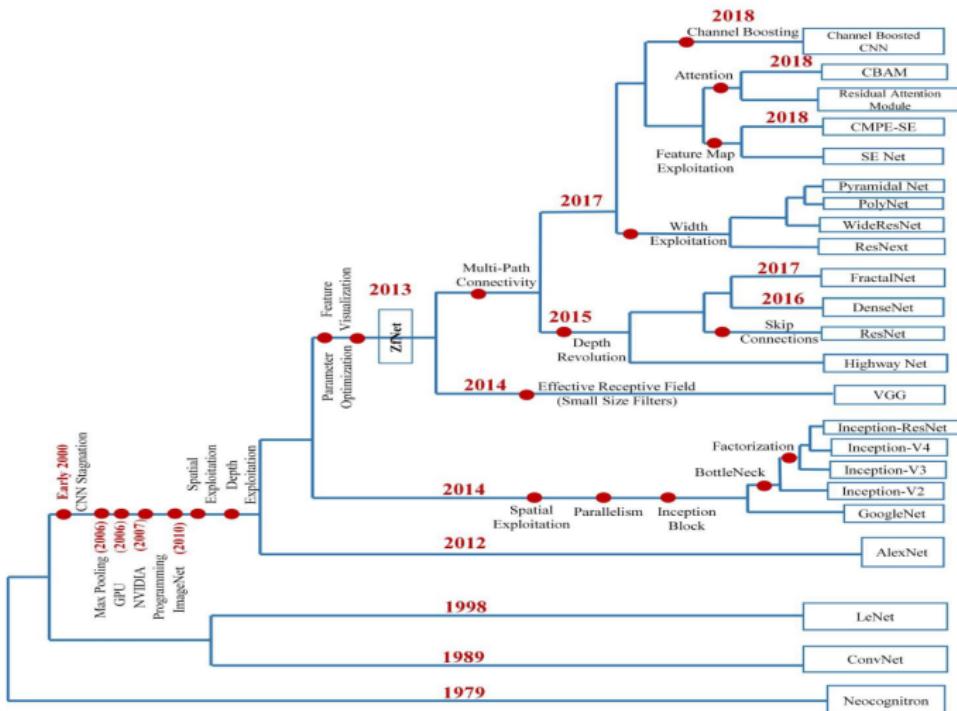
Layer	Shape	Size	Parameter
Input layer	(32,32,3)	3072	0
CONV-1 (f=5, S=1, 6 filters)			
POOL-1			
CONV-1 (f=5, S=1, 16 filters)			
POOL-2			
FC3			
FC4			
Softmax			

Parameters in CNN

Layer	Shape	Size	Parameter
Input layer	(32,32,3)	3072	0
CONV-1 (f=5, S=1, 6 filters)	(28,28,6)	4704	456
POOL-1	(14,14,6)	1176	0
CONV-1 (f=5, S=1, 16 filters)	(10,10,16)	1600	2416
POOL-2	(5,5,16)	400	0
FC3	(120,1)	120	48120
FC4	(84,1)	84	10164
Softmax	(10,1)	10	850

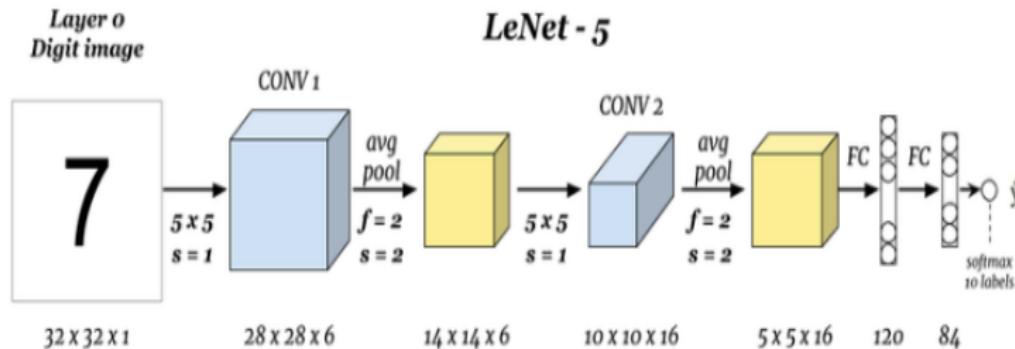
Total number of learnable parameters: 62006

Evolution of image classifiers



LeNet

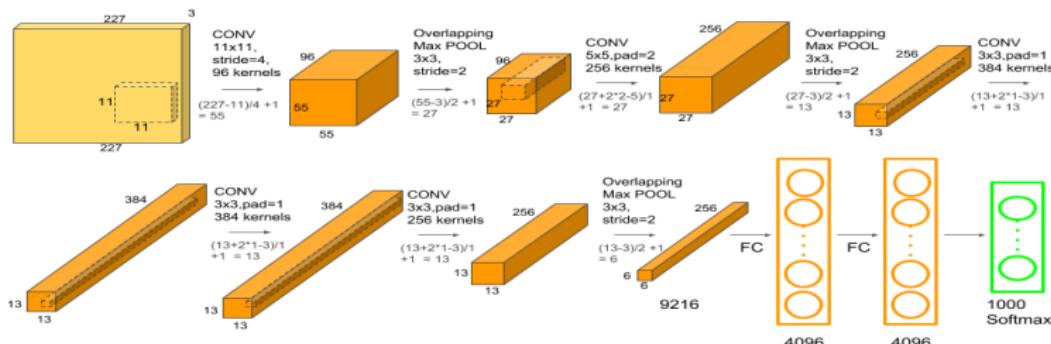
- Developed in 1998 and used for pincode digit recognition in postal service
- Around 60k parameters
- Sigmoid and Tanh were used



Lecun et al, 1998

AlexNet

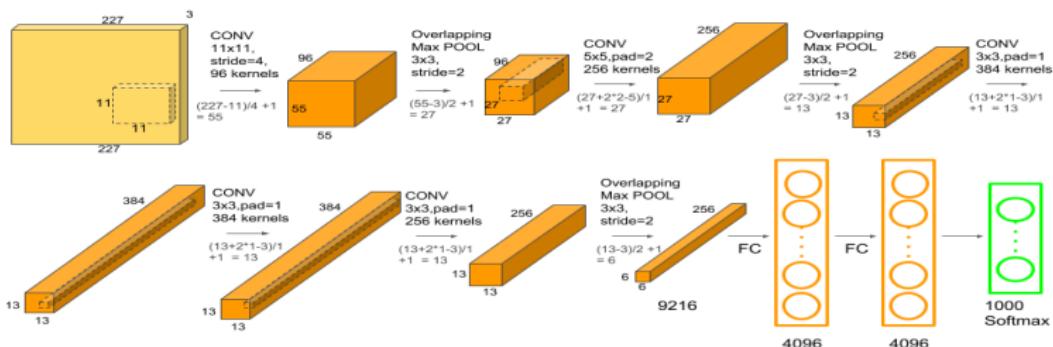
- Developed in 2012, similar to LeNet but with more layers
- 60 Million parameters
- Relu activation and local response normalization (LRN) is used



Alex K et al, 2012

VGG16

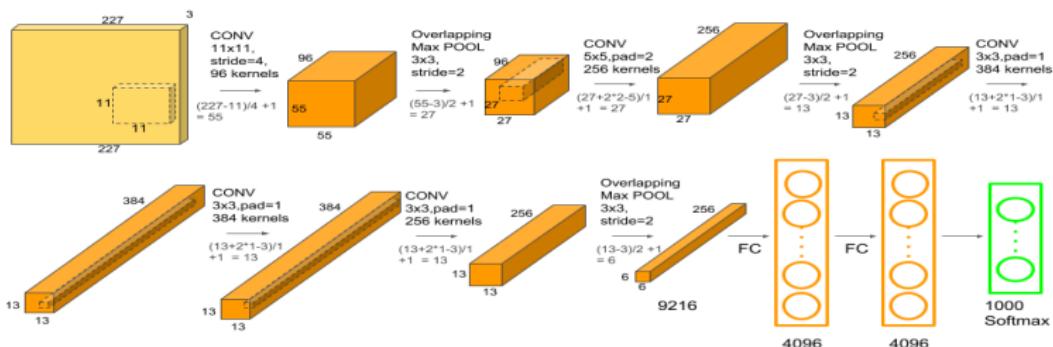
- Developed in 2015, CONVNET with 3 x 3 filters
- 16 layers with 138 Million parameters
- Number of Filters: 64 → 128 → 256 → 512
- Error rate is 17% compared to 8% in AlexNet



Simonyan and zisserman, 2015

ResNet

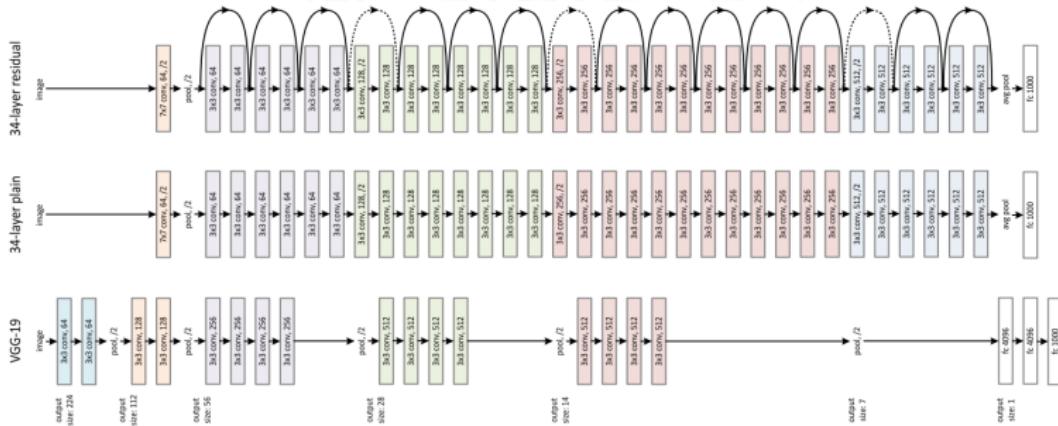
- Developed in 2015, CONVNET with 3 x 3 filters
- 16 layers with 138 Million parameters
- Number of Filters: 64 → 128 → 256 → 512
- Error rate is 17% compared to 8% in AlexNet



Simonyan and zisserman, 2015

Residual Network

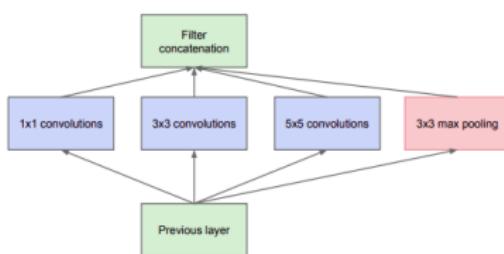
- Very deep NN has problem of vanishing and exploding gradients
 - Skip connections are used which helps in training very deep layers



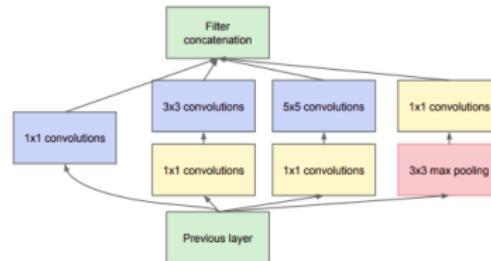
He et al, 2015

GoogLeNet

- Homage to Yann LeCuns pioneering LeNet-5 network
- 1×1 convolution to shrink larger volume to smaller volume
- Inception module to use different filter in same layer



(a) Inception module, naïve version

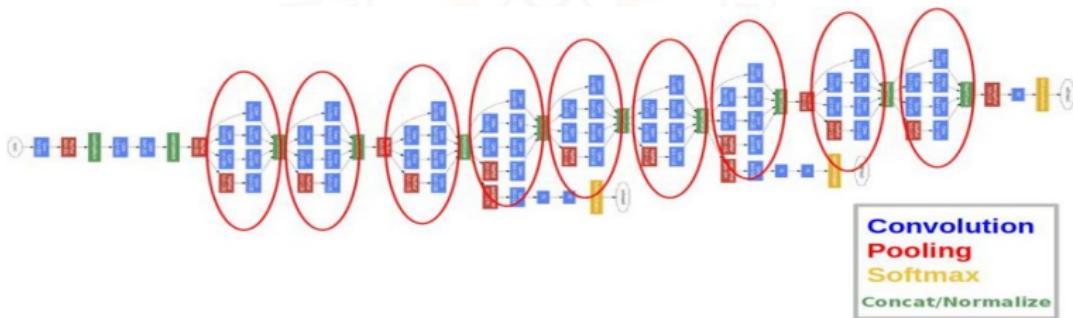


(b) Inception module with dimension reductions

Christian et al, 2015

GoogLeNet

- 22 layers deep network with inception module
- Quality gain at modest increase in computational requirement

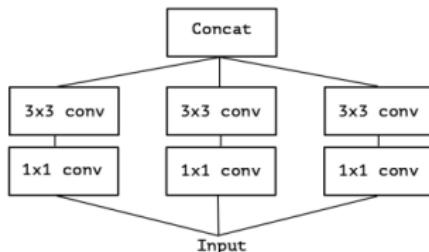


Christian et al, 2015

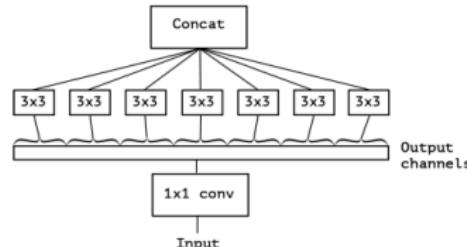
Xception Network

- An extreme version of inception module
- Deepwise seperable convolutions in NN architecture
- Similar parameters compared to inception V3 with gain in classification performance

A simplified Inception module.



An “extreme” version of our Inception module, with one spatial convolution per output channel of the 1×1 convolution.



Choller, 2017

Transfer learning

- CNN are basically feature extractors which learn to perform image classification
- A model trained on ImageNet Challenge have learnt to extract features from a wide range of images.
- **Can we transfer this knowledge to solve some other problems as well?**

Transfer learning is the practice of reusing the skills learnt from solving one problem to solve a new related problem

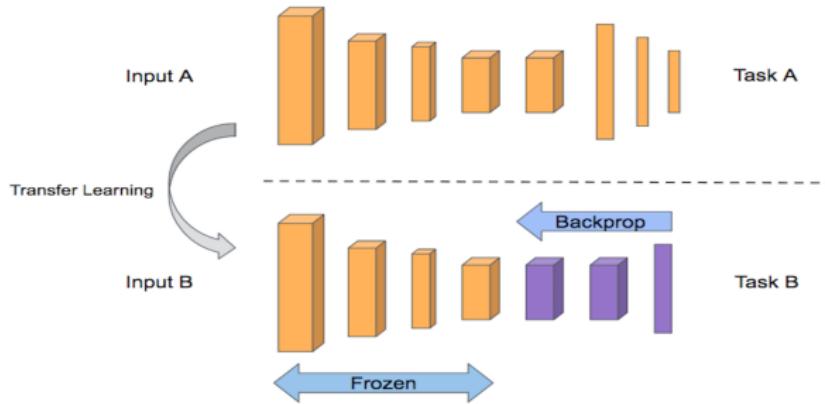
Ex: Bicycle and two-wheeler

Main reasons for Transfer learning

- Less data for particular problem: Data abundance in one task and data crunch in other related task
Eg: Model for driverless car in Indian roads
- No computational power (enough data available for training)
- For most computer vision problems, pre-trained models such as AlexNet, VGGNet, GoogleNet etc works better.
Eg: Coal mine and elephant feature extraction.
- Last few layers are only used for discrimination and there are lot of learning happened in previous layer

Ways of Transfer learning

- Freeze weights of initial few layers and train few later layers
- Retain the entire network of all weights
- If we have lot of data, then we use open source architecture and train the whole model



PREDICTION: IITM'S CRICKET AND CODE CHALLENGE



Problem Statement:

- For upcoming IPL T20 matches, participants should submit their code to predict the score at the end of 6 overs in each of the innings.
- Given: Ball-by-ball data from past 1600 innings
- Expected output: For each innings of each match, the score at the end of the 6th over.
- Input: Team batting, team bowling, name of the batsman and bowler during the six overs

PREDICTION: IITM'S CRICKET AND CODE CHALLENGE



Data preparation:

- From the 1600 matches, we need to separate IPL matches with following features as input data
 - Batting team name (CSK-1, KKR-2, SRH-3 etc)
 - Bowling team name (CSK-1, KKR-2, SRH-3 etc)
 - No of wickets (1,2,3...)
 - No of bowlers used during six overs (1,2,3...)
 - Name of the ground (MAC-1, Wankade-2, kotla-3 etc.)
 - Runs scored during first six overs (30,40,65...)
 - Batting first/chasing (0/1)

PREDICTION: IITM'S CRICKET AND CODE CHALLENGE

Sample data file:

Batting team	Bowling team	Ground	Number of wickets	Number of bowlers	Batting first	Runs scored
1	2	1	3	2	1	45
2	1	2	1	3	0	54
4	5	4	3	2	1	36
5	4	3	2	4	0	70

PREDICTION: IITM'S CRICKET AND CODE CHALLENGE

Models developed:

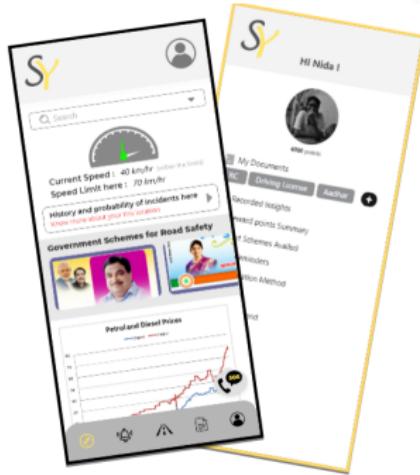
- Team specific model
- Ground specific model
- General model

Based on the cricket knowledge, we selected which model to use before each match.

We are among the top 15 teams in the list of 500 teams

Rather than type of model, features (input) plays an vital role in prediction

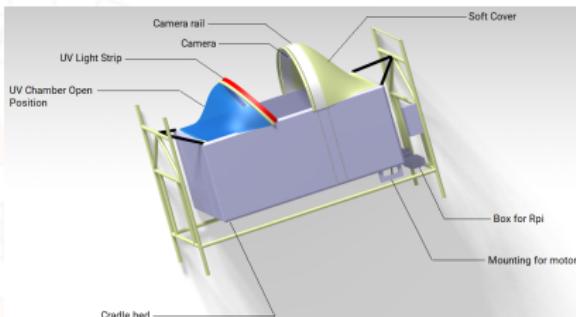
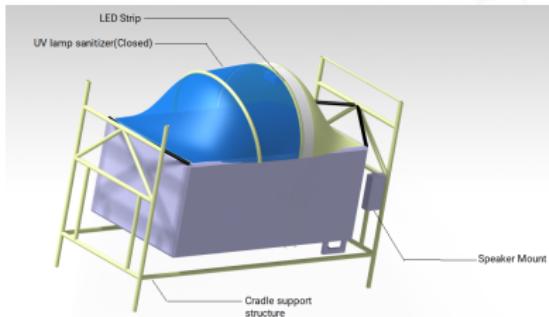
CLASSIFICATION: ROAD SAFETY IDEATION CHALLENGE



- ML based warning system for potholes
- Advanced warning system about blind spots, speed limits and upcoming traffic lights
- Checking the adherence of rider
- Collection of data for insurance claim

Second prize has been won by the team for this ideation with cash rewards of 35,000 from GI council

PRODUCT DESIGN: SMART INFANT MANAGEMENT SYSTEM



Features:

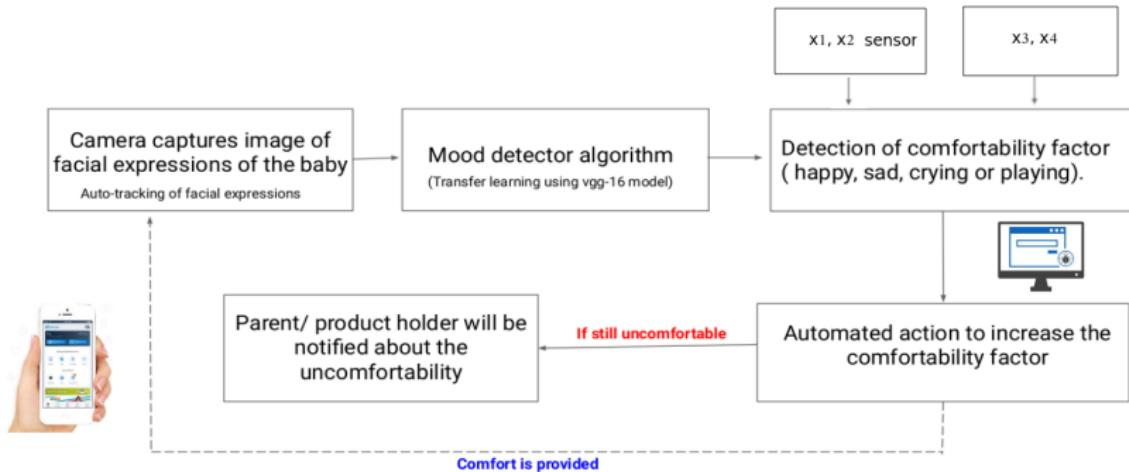
- UV sanitisation facility for the cardle
- ML based emotion classification
- Web and app control of cardle
- Medical data storage for analytics

Target users:

- Infant wards
- Dual working parents

PRODUCT DESIGN: SMART INFANT MANAGEMENT SYSTEM

Emotion classification overview:



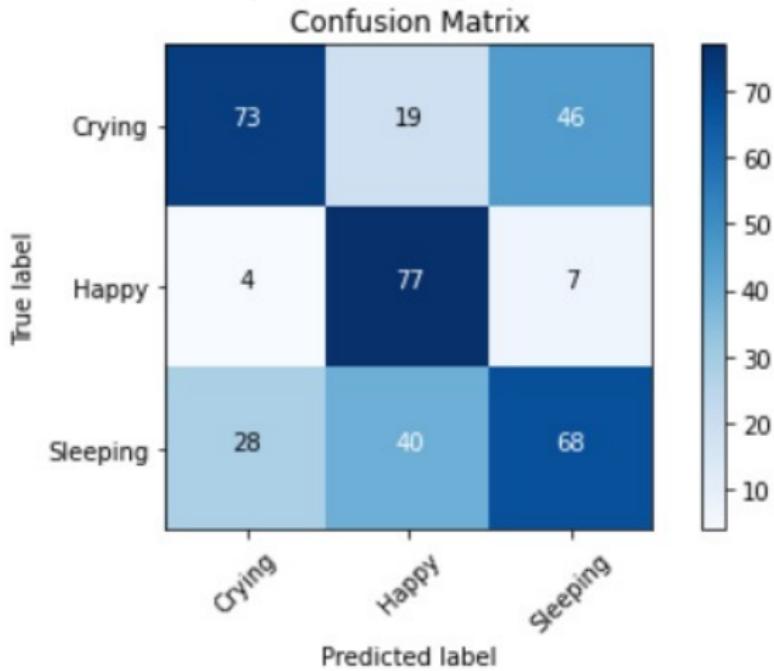
PRODUCT DESIGN: SMART INFANT MANAGEMENT SYSTEM

Dataset collection:

- Three categories of emotion
 - Happy
 - Sleeping
 - Crying
- Datasets are collected from google, instagram videos, youtube videos and directly approaching parents
- Dataset used
 - Training set (4.5k images)
 - Validation set (750 images)
 - Test set (362 images)

PRODUCT DESIGN: SMART INFANT MANAGEMENT SYSTEM

Confusion matrix for the model X:



TEAM MEMBER PACKAGES



SAMSUNG

Wiingy
The Technology School

incedo

$$60 + 46 + 44 + 40 + 20 + 22 + 8$$